

Ch.6 SDN(2)

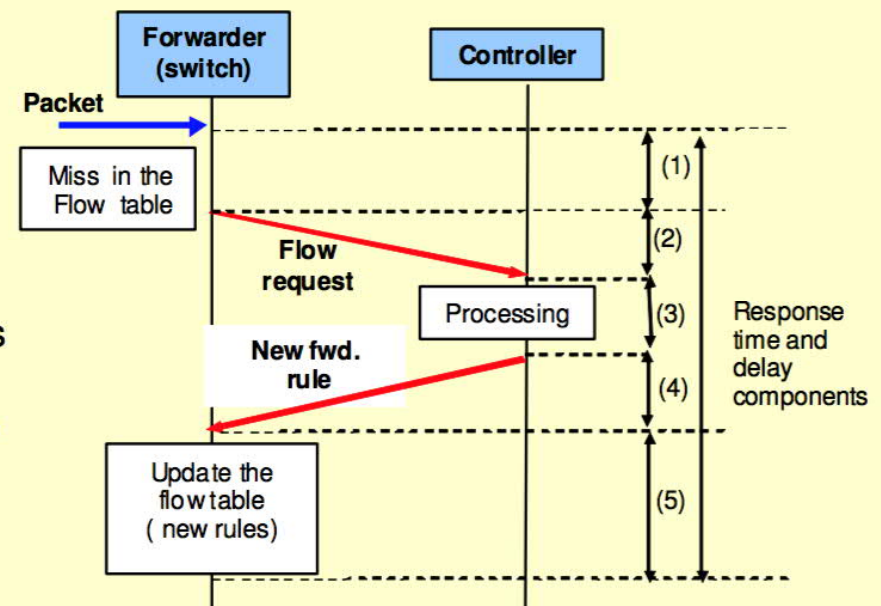
Arch & APP

Outline

- Arch
 - Scalability
 - Programming
- App
 - Security

■ SDN Problem: new flows setup - response time

- Signaling overhead components
 - Switch (CPU, mem, ..) - (1), (5)
 - Controller (CPU, mem, ..) - (3)
 - Transport through network – (2), (4)
- Transport: (2), (4) → place controller closer to the switch
- Switch
 - OpenVSwitch: install tens of (10^{**3}) flows/s with < 1 ms latency
 - HW switches: install few (10^{**3}) with 10ms latency
 - Weak mgmt CPU
 - Low speed communication CPU-switching chipset
 - **Hope for faster switches**
- Frequent events can stress
 - Controller resources
 - Control channel
 - Switch



On demand flow setup

- **Solutions:**
 - **Direct solutions**
 - Increase controller processing power
 - Increase switch processing power
 - **Aggregation of rules**
 - Proactive installation of rules
 - Problems: no host mobility support, not enough memory in switches
 - **Delegate more responsibilities to the data plane**
 - to switch control plane [e.g. Diffane, DevoFlow]
 - **Distributed controllers**
 - Flat structure multiple controllers [e.g. ONIX]
 - Recursive controller design [e.g. Xbar]
 - Hierarchical controller design [e.g. Kandoo]



3. Control Plane Scalability in SDN



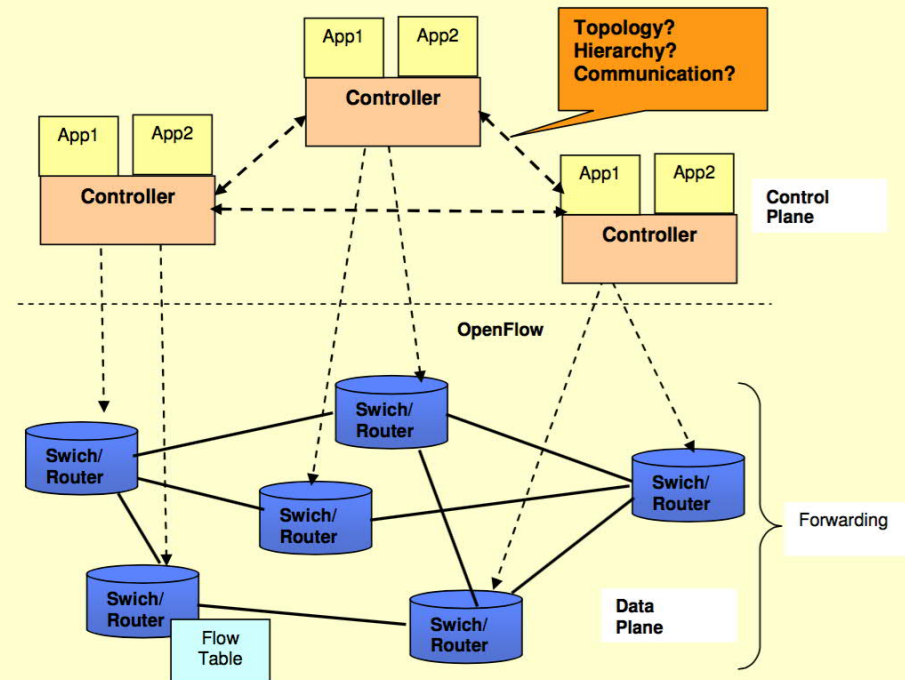
- **SDN Scalability-related solutions**
- **Increase the controller processing power**
- *Source: A. Tootoonchian et al., "On Controller Performance in Software-Defined Networks," Proc. USENIX Hot-ICE '12, 2012, pp. 10–10.*
 - Multicore systems for higher level of parallelism
 - And improved IO performance
 - Simple modifications to the NOX controller → performance increase more than 10 times an order of magnitude on a single core. (w.r.t 30000 flows/sec)
 - A single controller can support larger networks, (if the controller channel has enough bandwidth and acceptable latency)



3. Control Plane Scalability in SDN



- **SDN Scalability issues**
 - **Multiple controller solution**
 - Control Plane and Data Plane – independent; they can have different topologies





3. Control Plane Scalability in SDN



- **Multiple controllers- (cont'd)**
- **Approaches:**
 - **Flat organization**
 - Distribution of the CPI while maintaining a logically centralized using a distributed file system, a distributed hash table and a pre-computation of all possible combinations respectively.
 - however they impose a strong requirement: a consistent network-wide view in all the controllers.
 - Examples: HyperFlow, Onix and Devolved controllers
 - **Hierarchical organization**
 - Hierarchical distribution of the controllers based on two layers:
 - (i) the bottom layer, a group of controllers with no interconnection, and no knowledge of the network-wide state,
 - (ii) the top layer, a logically centralized controller that maintains the network wide state.
 - Example: Kandoo



3. Control Plane Scalability in SDN



- **Solution proposals examples**

- **HyperFlow**

- *Source: A. Tootoonchian et.al., "Hyperflow: A Distributed Control Plane for OpenFlow," Proc. 2010 INMConf., 2010.*
- Among the first distributed (event-based) Control Plane for OpenFlow
- **Logically centralized**
 - All the controllers share the same consistent network-wide view and locally serve requests without actively contacting any remote node, thus minimizing the flow setup times.
- **Physically distributed**: scalable while keeping the network control centralization benefits
- By passively synchronizing network-wide views of OpenFlow controllers, it localizes decision making to individual controllers, thus minimizing the CPI response time to data plane requests
- It is resilient to network partitioning and component failures
- **It enables interconnecting independently managed OpenFlow networks** - - an essential feature



3. Control Plane Scalability in SDN



- **Solution proposals examples**
- **HyperFlow (cont'd)**
 - **Design Principles**
 - A HyperFlow (HF) -based network is composed by
 - **OpenFlow switches/forwarders**
 - **NOX controllers** (decision elements) each running an instance of the HF controller application
 - **Event propagation system** for cross-controller communication.
 - All the controllers have a consistent network-wide view and run as if they are controlling the whole network.
 - They **all run the exact same controller software and set of applications.**
 - Each switch is connected to the closest controller
 - Controller failure → affected switches must be reconfigured to connect to an active nearby controller
 - **Each controller**
 - directly manages a subset of switches connected to it
 - indirectly programs or queries the rest (via comm.. with other controllers).




3. Control Plane Scalability in SDN



- Solution proposals examples
- HyperFlow (cont'd)
 - High level view


Event propagation system for cross-controller communication
Solution:
publish/subscribe system



Outline

- Arch
 - Scalability
 - Programming
- App
 - Security

Frenetic: A Programming Language for Software Defined Networks



Jennifer Rexford
Princeton University



<http://www.frenetic-lang.org/>


Joint work with Nate Foster, Dave Walker, Rob Harrison, Michael Freedman, Chris Monsanto, Mark Reitblatt, and Alec Story

Network Programming is Hard

- Programming network equipment is hard
 - Complex software by equipment vendors
 - Complex configuration by network administrators
- Expensive and error prone
 - Network outages and security vulnerabilities
 - Slow introduction of new features
- SDN gives us a chance to get this right!
 - Rethink abstractions for network programming


Programming Software Defined Networks

- OpenFlow already helps a lot
 - Network-wide view at controller
 - Direct control over data plane
- The APIs do not make it easy
 - Limited controller visibility
 - No support for composition
 - Asynchronous events
- Frenetic simplifies the programmer's life
 - A language that raises the level of abstraction
 - A run-time system that handles the gory details



Limited Controller Visibility

- Example: MAC-learning switch
 - Learn about new source MAC addresses
 - Forward to known destination MAC addresses
- Controller program is more complex than it seems
 - Cannot install destination-based forwarding rules
 - ... without keeping controller from learning new sources




- Solution: rules on <inport, src MAC, dst MAC>

Must think about *reading* and *writing* at the same time.

Composition: Simple Repeater

Simple Repeater

```
def switch_join(switch):  
    # Repeat Port 1 to Port 2  
    p1 = {in_port: 1}  
    a1 = [forward(2)]  
    install(switch, p1, DEFAULT, a1)  
  
    # Repeat Port 2 to Port 1  
    p2 = {in_port: 2}  
    a2 = [forward(1)]  
    install(switch, p2, DEFAULT, a2)
```




When a switch joins the network, install two forwarding rules.

Composition: Web Traffic Monitor

Monitor “port 80” traffic

```
def switch_join(switch):  
    # Web traffic from Internet  
    p = {inport: 2, tp_src: 80}  
    install(switch, p, DEFAULT, [])  
    query_stats(switch, p)  
  
def stats_in(switch, p, bytes, ...)  
    print bytes  
    sleep(30)  
    query_stats(switch, p)
```



When a switch joins the network, install one monitoring rule.

Composition: Repeater + Monitor


Repeater + **Monitor**

```
def switch_join(switch):  
    pat1 = {inport: 1}  
    pat2 = {inport: 2}  
    pat2web = {in_port: 2, tp_src: 80}  
    install(switch, pat1, DEFAULT, None, [forward(2)])  
    install(switch, pat2web, HIGH, None, [forward(1)])  
    install(switch, pat2, DEFAULT, None, [forward(1)])  
    query_stats(switch, pat2web)  
  
def stats_in(switch, xid, pattern, packets, bytes):  
    print bytes  
    sleep(30)  
    query_stats(switch, pattern)
```

Must think about both tasks at the same time.

Asynchrony: Switch-Controller Delays

- Common OpenFlow programming idiom
 - First packet of a flow goes to the controller
 - Controller installs rules to handle remaining packets



- What if more packets arrive before rules installed?
 - Multiple packets of a flow reach the controller
- What if rules along a path installed out of order?
 - Packets reach intermediate switch before rules do

Must think about all possible event orderings.

Wouldn't It Be Nice if You Could...

- Separate reading from writing
 - Reading: specify queries on network state
 - Writing: specify forwarding policies
- Compose multiple tasks
 - Write each task once, and combine with others
- Prevent race conditions
 - Automatically apply forwarding policy to extra packets

This is what Frenetic does!

Our Solution: Frenetic Language

- Reads: query network state
 - Queries can see any packets
 - Queries do not affect forwarding
 - Language designed to keep packets in data plane
- Writes: specify a forwarding policy
 - Policy separate from mechanism for installing rules
 - Streams of packets, topology changes, statistics, etc.
 - Library to transform, split, merge, and filter streams
- Current implementation
 - A collection of Python libraries on top of NOX

Example: Repeater + Monitor

```
# Static repeating between ports 1 and 2
def repeater():
    rules=[Rule(inport: 1, [forward(2)]),
           Rule(inport: 2, [forward(1)])]
    register(rules)
```

Repeater

```
# Monitoring Web traffic
def web_monitor():
    q = (Select(bytes) *
         Where(inport: 2 & tp_src: 80) *
         Every(30))
    q >> Print()
```

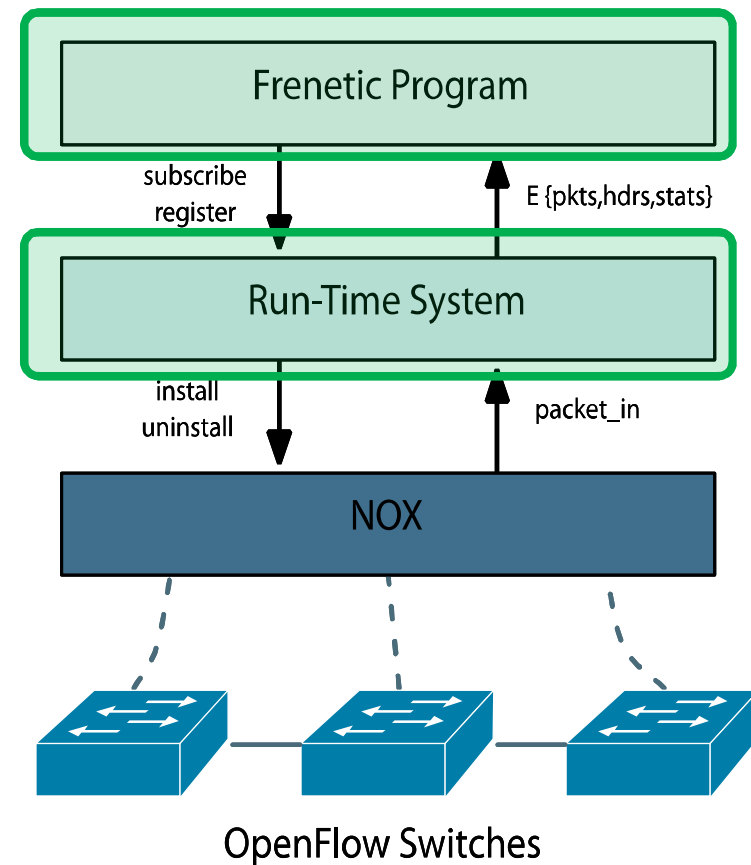
Monitor

Repeater + Monitor

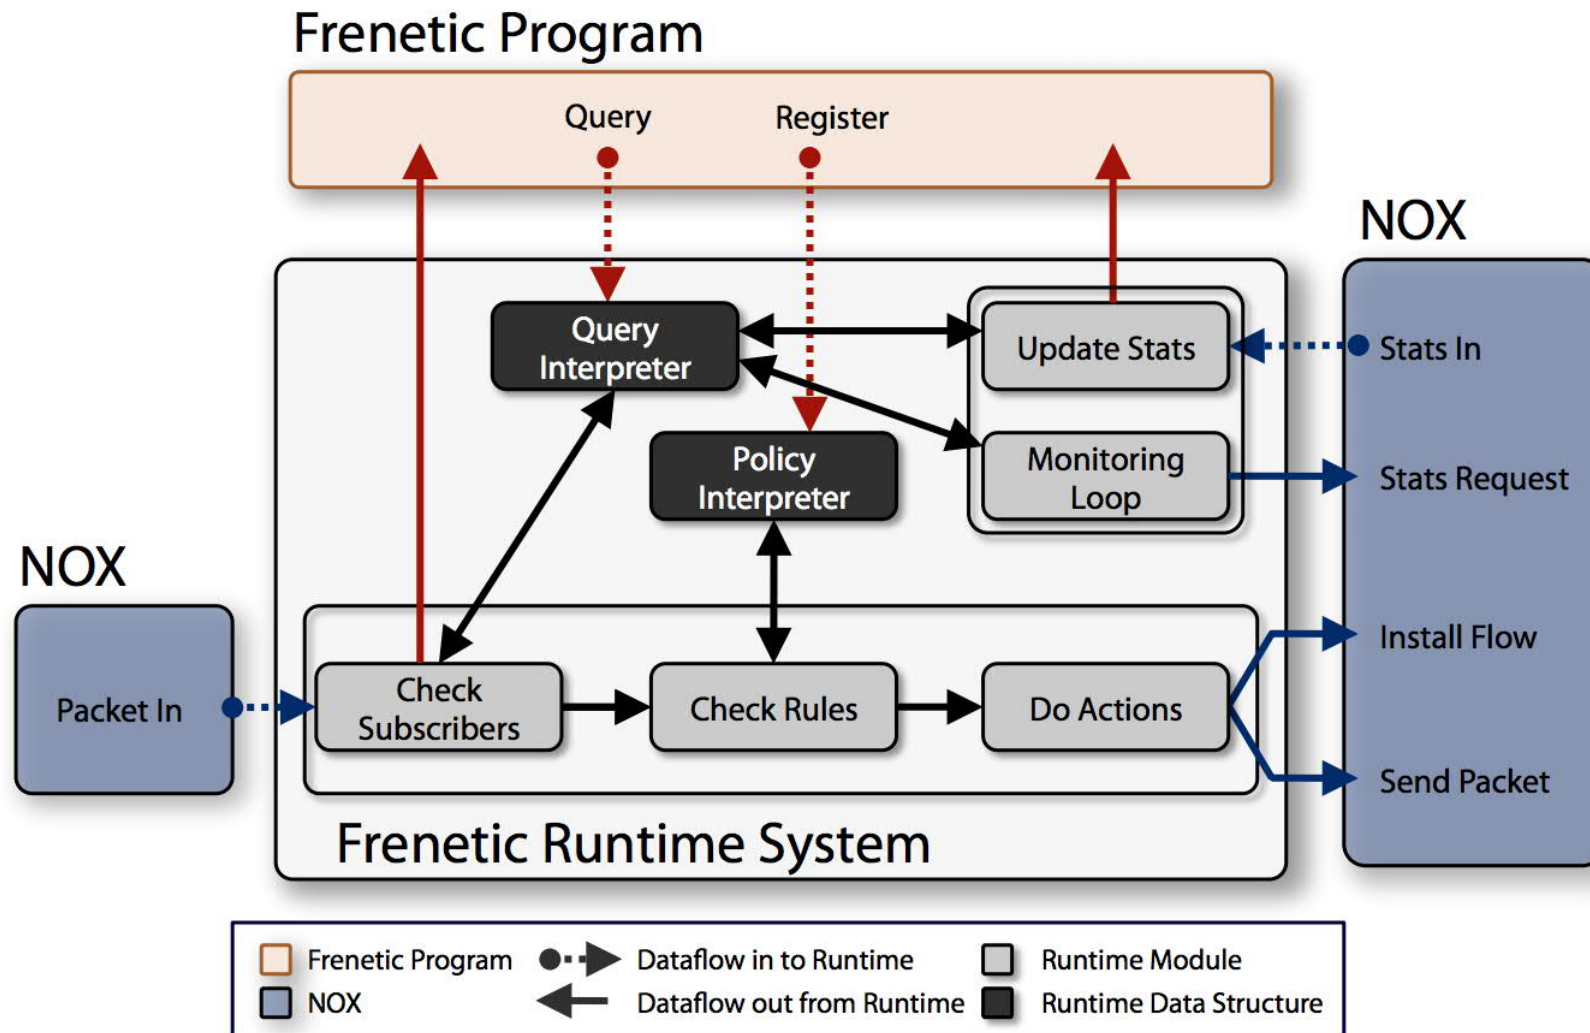
```
# Composition of two separate modules
def main():
    repeater()
    web_monitor()
```

Frenetic System Overview

- High-level language
 - Query language
 - Composition of forwarding policies
- Run-time system
 - Interprets queries and policies
 - Installs rules and tracks statistics
 - Handles asynchronous events



Run-time Activities

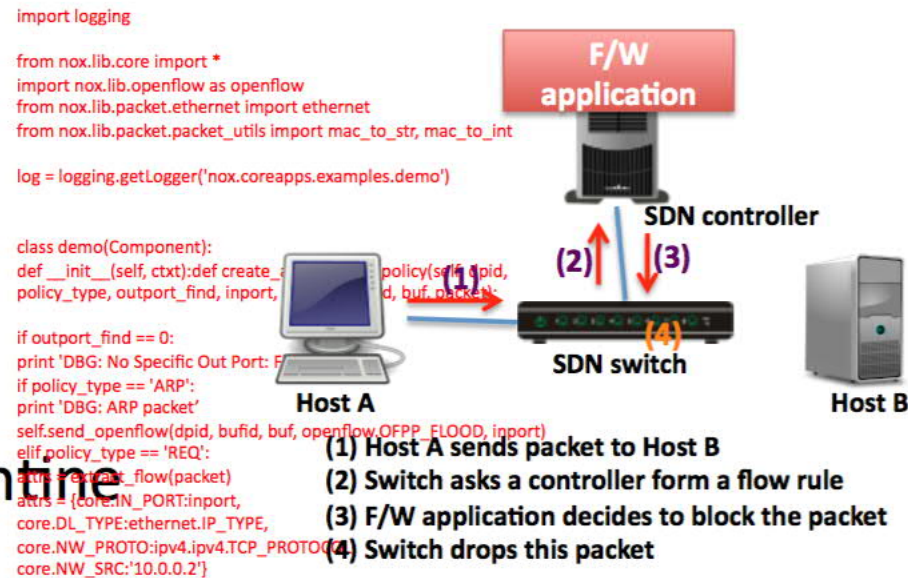


Outline

- Arch
 - Scalability
 - Programming
- App
 - Security

Exemplar SDN Security Apps

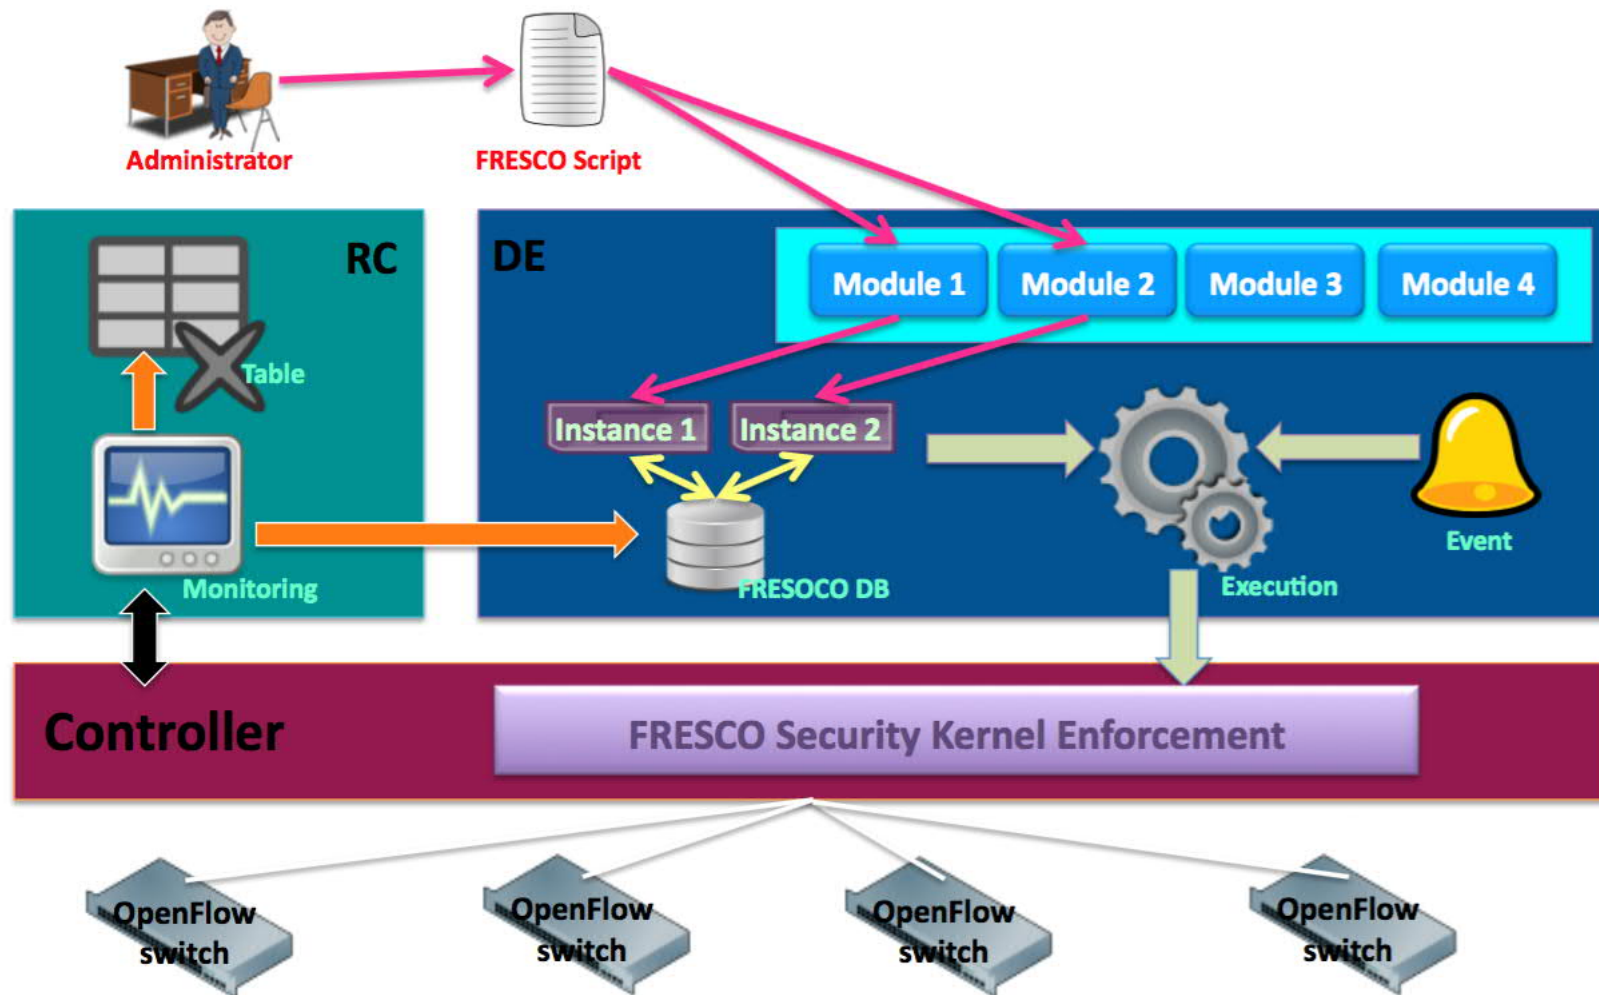
- Security functions can be applications of SDN
 - Firewall
 - DDoS detection
 - Scan detection
 - Reflector net
 - Tarpit
 - Dynamic quarantine
 - and more...



FRESCO

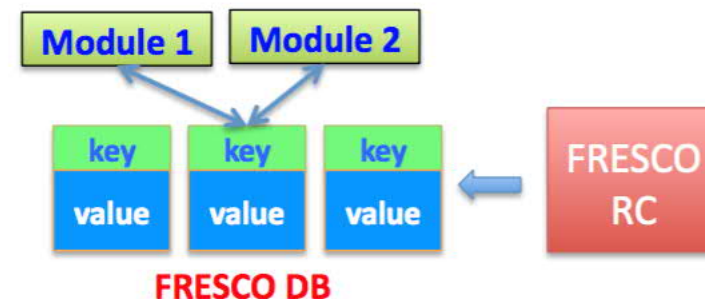

- FRESCO is a new frame work that
 - Provides a new development environment for security applications
 - Effectively manages shared resources among security applications
 - Simplifies deployment of security policies
 - provides a set of 7 new intelligent security action primitives
 - E.g., block, deny, allow, redirect, and quarantine

Operational Scenario



Development Environment

- FRESCO Module
 - Basic operation unit
- FRESCO DB
 - Simple database
 - (key,value) pairs
- FRESCO script
 - Define interfaces
 - Connect multiple modules



Development Environment

- FRESCO script

- Format

- **Instance name (# of input) (# of output)**
 - **type**: class of this module
 - **input**: input for this module
 - **output**: output of this module
 - **parameter**: define some variables
 - **event**: trigger a module
 - **action**: conduct this action

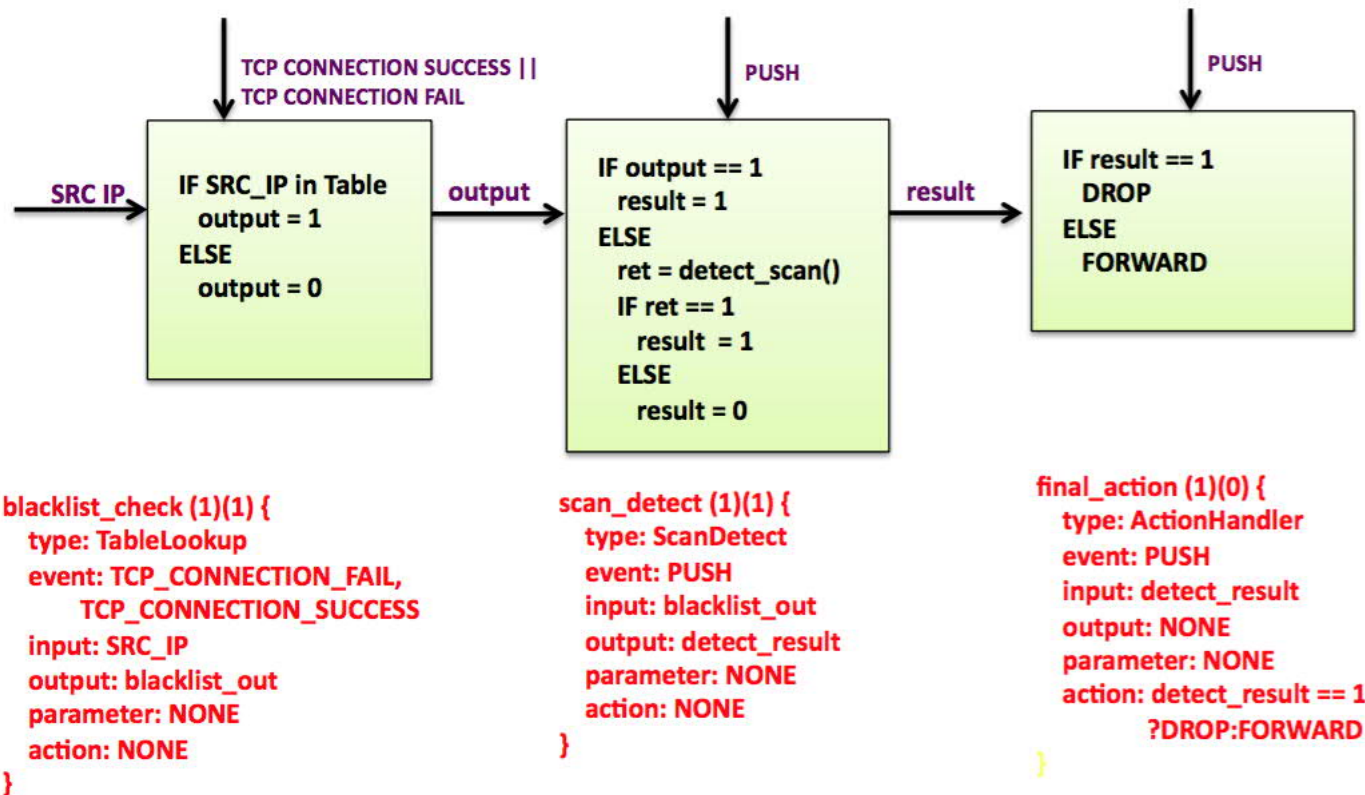
```
port_comparator (1)(1) {  
  type: Comparator  
  event: PUSH  
  input: destination_port  
  output: comparison_result  
  parameter: 80  
  action: -  
}
```

Inspired by Click Modular Router

Robert Morris, Eddie Kohler, John Jannotti, and M. Frans Kaashoek. Proceedings of SOSP '99

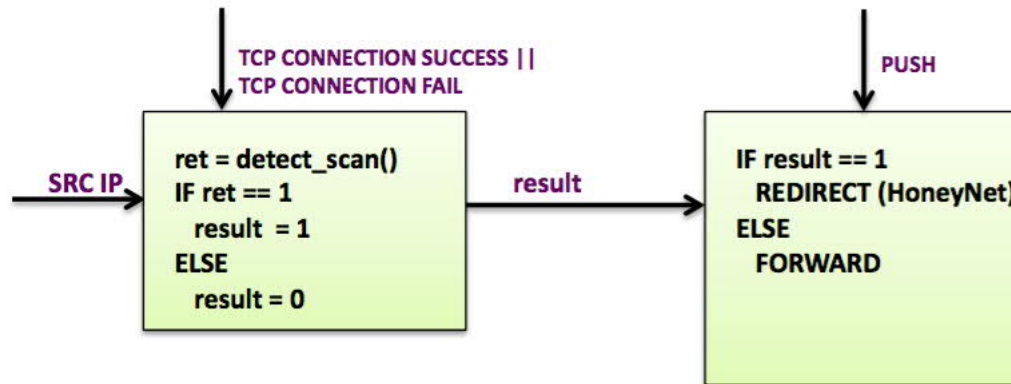
Example: Scan Detection

- Steps
 - Check blacklist → Threshold based scan detection → Drop or Forward



Example: Reflector Net

- Confuse network scan attackers
- Steps
 - Threshold based scan detection → Reflect or Forward



```
scan_detect (1)(1) {  
  type: ScanDetect  
  event: TCP_CONNECTION_FAIL,  
         TCP_CONNECTION_SUCCESS  
  input: blacklist_out  
  output: detect_result  
  parameter: NONE  
  action: NONE  
}
```

```
final_action (1)(0) {  
  type: ActionHandler  
  event: PUSH  
  input: detect_result  
  output: NONE  
  parameter: NONE  
  action: detect_result == ?REDIRECT(10.0.0.3):FORWARD  
}
```


Example: Reflector Net

- Test result

